

YAP470

Fatma Zehra Paksoy and Samet Emre Demir

September 2024

1 Abstract

Estimating vehicle prices is tricky since many factors come into play, such as age, brand, mileage covered, and fuel type. This research investigates this challenge by utilizing regression-based machine learning models to effectively estimate the prices of vehicles. The "Car Price Prediction" dataset sourced from Kaggle contains 301 observations and 9 key features that were used in the training, validation, and evaluation of the model.

Among the compared models, the best performance was obtained with the Random Forest model, which had a Mean Squared Error (MSE) of 0.1284 and an R^2 score of 0.8716. Support Vector Regression and CatBoost Regression also produced competitive results; however, Random Forest outperformed these methods in both accuracy and error metrics. This work emphasizes how preprocessing steps-features normalization, outlier detection, and categorical encoding-can enhance the performance of the model.

The results are in agreement with the previous literature and point out the robustness of ensemble methods in modeling complicated relationships in vehicle pricing. This framework offers a robust solution for stakeholders in the automotive industry to develop strategic pricing policies in the used vehicle market.

2 Introduction

Estimating the exact price of a vehicle is quite challenging owing to ever-changing market conditions and a wide range of influencing factors on the price, including but not limited to age, brand, mileage, and fuel type. The topic has attracted considerable interest in recent years from both academia and the automotive industry and thus prompted various studies using both traditional statistical methods and modern machine learning approaches.

For example, Erdem and Şentürk (2009) conducted hedonic regression on the prices of used cars in Turkey, studying specific features related to, for example, engine type and color. Li and Lin (2018) illustrate the potential for deep learning methods, such as Multilayer Perceptrons, to achieve high accuracy in prediction. Most of these studies are either limited in handling complex datasets with high

feature variance or are unable to make extensive comparisons among diverse machine learning and deep learning algorithms.

This paper investigates the use of a Kaggle publicly available dataset called "Car Price Prediction" containing 301 observations with 9 features, including some critical attributes like vehicle model, brand, mileage, and fuel type. A holistic approach was followed by utilizing a few regression-based machine learning models, including Random Forest, Support Vector Regression, and Lasso. Also, robust preprocessing like normalization, outlier removal, and categorical encoding were used. Amongst all these, the random forest with feature normalization topped others with an MSE of 0.1284 and R^2 score of 0.8716.

By combining different preprocessing methodologies with a wide variety of machine learning models, this paper tries to predict car prices more accurately and reliably. This work gives very useful insights and a strong framework for the decision-makers in the automotive market to support the price determination of second-hand vehicles.

3 Related Work

Used car price prediction has now become a very important thrust area of research, comprising state-of-the-art machine learning and deep learning methods. In this section, prior works are grouped into the following: data utilization and preprocessing, machine learning methodologies, deep learning approaches, and regional studies. This review further helps in identifying the lacunae in the existing literature and puts our contributions in perspective.

3.1 Data Utilization and Preprocessing

Effective data preprocessing is a must for any accurate price prediction. Erdem and Şentürk (2009) [1] used hedonic regression to analyze the Turkish second-hand car market based on features such as engine type and color. While foundational, their work lacked the complexity introduced by modern ML techniques.

Şahin (2023) [7] used Kaggle's dataset and highlighted preprocessing steps such as cleaning and visualization of data in order to make it ready for ML models. Gülmez and Kulluk (2023) [8] applied Apache Spark for preprocessing large datasets, using algorithms like Random Forest to identify key features. Similarly, Yılmaz (2023) [9] incorporated PCA and feature engineering to optimize datasets, enhancing model performance. These studies underscore the importance of robust preprocessing in achieving high prediction accuracy. Our work builds on these by implementing feature normalization, outlier removal, and categorical encoding to refine the dataset.

3.2 Machine Learning Techniques

Machine learning has been a dominant approach in vehicle price prediction research. Jin (2021) [3] illustrated the superiority of Random Forest over other models like linear and polynomial regression and support vector machines for handling nonlinear relationships. Similarly, Pal et al. (2018) [4] also mentioned the robustness of Random Forest in managing complex datasets.

Dere and Kasapbaşı (2023) [13] compared Random Forest and Tree Ensemble algorithms using datasets gathered from web scraping. Their study very well demonstrated the effectiveness of automated data acquisition in increasing the accuracy with the help of ML models. However, most of these works are restricted to a comparative study of a few chosen algorithms. Our research extends this study by comparing multiple ML models, including Random Forest, Support Vector Regression, and CatBoost, to identify the most effective approach.

3.3 Deep Learning Approaches

Deep learning has reshaped the horizon of price prediction with the ability to model complicated relations. Li and Lin (2018) [2], illustrated an effectiveness check on multilayer perceptron networks; thus, attaining substantially high accuracy over conventional regression models. Nevertheless, their study didn't draw a comparison with several machine learning ensembles such as Random Forest.

Alnajim et al. (2024) [6] used deep neural networks to predict car prices within the Saudi Arabian market and attained high accuracy with localized datasets. Erol (2024) [12] integrated DL models with Flask-based deployment to enable real-time predictions. While DL methods offer scalability and adaptability, they often require significant computational resources and are less interpretable than ML algorithms. Our work strikes a balance in focusing on ML techniques which have demonstrated robust performance without the attendant computational overhead of the DL models.

3.4 Regional and Specialized Studies

Various studies have designed their methods for specific regions or markets. For example, Namlı et al. (2019) [10] applied ML methods such as support vector machines and artificial neural networks to the Turkish second-hand car market and showed that ML methods outperform traditional regression models on regional datasets. Arefin (2021) [11] developed a price predictor for Tesla's second-hand market using Gradient Boosted Decision Trees. Although these domain-specific methods are effective, they often do not generalize well to larger markets.

Our research fills these gaps by utilizing a dataset that is diverse in feature variety, applicable to the Turkish second-hand car market. The aim is to develop a model that balances generalizability with region-specific insights.

3.5 Summary and Contribution

The literature review reflects a significant enhancement in the methodologies of price prediction, which ranges from traditional statistical approaches to advanced ML and DL techniques. Factors such as data preprocessing, algorithm selection, and region-specific customization remain key in determining model performance. Based on these elements, our research incorporates thorough steps in preprocessing, multiple ML models, and hyperparameter tuning to come up with a reliable framework for house price prediction. By integrating the strengths of prior works, our approach contributes a precise and generalizable model for predicting vehicle prices in the Turkish second-hand car market.

4 Data

4.1 Data Source

The dataset used in this project was obtained from Kaggle, titled "*Car Price Prediction*", contributed by a user named *hellbuoy*. It contains a total of 301 records, each characterized by 9 attributes:

- **car_name**: The car model name.
- **year**: Car manufacturing year.
- **selling_price**: Selling price of the used car (target variable).
- **present_price**: Current price at the showroom.
- **kms_driven**: The number of kilometers already traveled by the car.
- **fuel_type**: Type of fuel used by the car (categorical: Petrol, Diesel, or CNG).
- **seller_type**: Type of seller (categorical: Individual or Dealer).
- **transmission**: Type of transmission (categorical: Manual or Automatic).
- **owner**: Number of previous owners.

4.2 Preprocessing

A structured preprocessing pipeline was applied to prepare the dataset for machine learning. The following steps were included:

4.2.1 Handling Missing Values

The dataset was checked for missing values, and none were found.

4.2.2 Outlier Detection and Treatment

- Box plots were used to identify outliers in **kms_driven** and **selling_price**.
- Extreme values were capped using percentile-based methods to maintain consistency in the dataset.

4.2.3 Encoding Categorical Variables

Categorical attributes (**fuel_type**, **seller_type**, and **transmission**) were transformed into numerical values using OneHotEncoder.

4.2.4 Feature Scaling

Continuous features (**kms_driven**, **present_price**, and **year**) were standardized using StandardScaler to normalize their values.

4.2.5 Feature Selection

- The **car_name** attribute was dropped as it was irrelevant for predicting the target variable.
- Other attributes were retained due to their correlation with **selling_price**.

4.3 Dataset Split

The preprocessed dataset was divided into two subsets for model training and testing:

- **Training Set (70%)**: 210 records used for model training.
- **Test Set (30%)**: 91 records reserved for final model evaluation.

4.4 Post-Preprocessing Statistics

- **Total Records**: 301
- **Training Records**: 210
- **Test Records**: 91
- **Label Distribution**:
 - **Fuel Type**: Petrol (60%), Diesel (35%), CNG (5%).
 - **Seller Type**: Individual (80%), Dealer (20%).
 - **Transmission Type**: Manual (85%), Automatic (15%).




Figure 1: Fuel Type Distribution.




Figure 2: Seller Type Distribution.




Figure 3: Transmission Type Distribution.

4.5 Feature Importance

The importance of each feature was evaluated during training to identify the most influential variables in predicting the selling price:

- The most significant predictors were **present_price** and **year**.
- **Kms Driven** had a moderate influence on the prediction.
- Attributes like **fuel_type**, **seller_type**, and **transmission** had relatively lower impacts.

5 System Design

The system is designed to predict the selling price of a used car with the highest accuracy using regression-based machine learning models. Further, this section describes the design of the system, including the input-output pipeline, preprocessing steps, model configurations, and evaluation metrics.

5.1 Input and Output

Input:

- The system takes in features such as **Present Price**, **Year**, **Kms Driven**, **Fuel Type**, **Seller Type**, **Transmission**, and **Owner**.
- These features are normalized and encoded to be compatible with the machine learning models.

Output:

- The system predicts the **Selling Price** of a car as a continuous value representing the estimated price in monetary terms.

5.2 Data Processing Pipeline

The input data is preprocessed through the following steps before being used for training and prediction:

- **Data Cleaning:** Missing values were checked and none were found. Outliers in **Kms Driven** and **Selling Price** were treated using percentile-based methods.
- **Feature Transformation:** Categorical variables (**Fuel Type**, **Seller Type**, **Transmission**) were encoded using OneHotEncoder, and continuous variables (**Present Price**, **Year**, **Kms Driven**) were standardized using StandardScaler.
- **Dataset Splitting:** The dataset was split into 70% training and 30% testing subsets.

5.3 Machine Learning Models

Several regression-based machine learning algorithms were employed. Below are the models and their configurations:

1. Random Forest Regression

- **Hyperparameters:**
 - Number of estimators: 100
 - Max depth: 10
 - Min samples split: 2
- **Rationale:** Combines multiple decision trees to provide robust predictions and minimize overfitting.

2. Support Vector Regression (SVR)

- **Hyperparameters:**
 - Kernel: Radial Basis Function (RBF)
 - Regularization parameter (C): 1
 - Epsilon: 0.1
- **Rationale:** Effective for handling non-linear relationships in the data.

3. Lasso Regression

- **Hyperparameters:**
 - Alpha: 0.01
- **Rationale:** Applies L1 regularization to reduce overfitting by penalizing large coefficients.

4. CatBoost Regression

- **Hyperparameters:**

- Depth: 8
- Learning rate: 0.03
- Iterations: 500

- **Rationale:** Efficiently handles categorical features and improves boosting performance.

5.4 Hyperparameter Tuning

Hyperparameter optimization was performed using GridSearchCV for the following configurations:

- **Random Forest:** Number of estimators (50, 100, 150) and max depth (5, 10, 15).
- **SVR:** Kernel type (linear, RBF) and regularization parameter (C: 0.1, 1, 10).
- **Lasso Regression:** Alpha (0.001, 0.01, 0.1).
- **CatBoost Regression:** Depth (6, 8, 10) and learning rate (0.01, 0.03, 0.05).

5.5 Evaluation Metrics

The models were evaluated using the following metrics:

- **Mean Squared Error (MSE):** Measures the average squared difference between predicted and actual values. Lower values indicate better performance.
- **Coefficient of Determination (R^2):** Indicates the proportion of variance in the dependent variable that is predictable from the independent variables. Values closer to 1 signify better explanatory power.

5.6 System Workflow

1. **Input Features:** User enters details such as car year, mileage, fuel type, seller type, and transmission type.
2. **Preprocessing:** The input data is cleaned, encoded, and scaled according to the preprocessing pipeline.
3. **Model Training:** The system trains models on the training dataset and tunes hyperparameters using the validation set.

4. **Prediction:** The system predicts the selling price of the car based on the test data, outputs the predicted value, and displays model performance metrics for transparency.

5.7 Conclusion

This systematic design ensures a robust framework for predicting car prices, accommodating various machine learning techniques and preprocessing strategies to achieve optimal results.

6 Results

This section presents the evaluation results of the models applied to predict used car prices. The performance metrics chosen for comparison are Mean Squared Error (MSE) and R^2 score, which evaluate the accuracy and explanatory power of the models. The results are summarized in Table 1, followed by a detailed analysis and insights.

6.1 Model Performance Comparison

The following table outlines the performance of the regression models tested in this study:

Model	Configuration	MSE	R^2	F1 Score	Precision	Recall
Random Forest	n_estimators=100, max_depth=10	3.62e+06	0.954	0.950	0.950	0.95
Elastic Net	alpha=0.01, l1_ratio=0.5	8.78e+06	0.889	0.837	0.783	0.90
Ridge Regression	alpha=1.0	8.75e+06	0.889	0.837	0.783	0.90
Lasso Regression	alpha=0.01	1.05e+07	0.867	0.857	0.818	0.90
k-NN	n_neighbors=5	1.60e+07	0.797	0.900	0.900	0.90
Linear Regression	Default configuration	1.15e+07	0.854	0.974	1.000	0.95
SVR	kernel=RBF, C=1, epsilon=0.1	7.02e+07	0.111	0.833	0.714	1.00

Table 1: Performance metrics for various regression models. The **MSE** (Mean Squared Error) represents the average squared prediction error, **R^2** indicates the proportion of variance explained by each model, and additional metrics such as **F1 Score**, **Precision**, and **Recall** are included to evaluate classification performance. Random Forest achieved the best performance overall.

6.2 Detailed Analysis of Results

The performance of the models can be analyzed as follows:

- **Random Forest:** Achieved the best results with an MSE of 3.62e+06 and an R^2 score of 0.954. Its ability to handle non-linear relationships and feature interactions contributed to its superior performance.
- **Elastic Net:** Demonstrated strong performance with an MSE of 8.78e+06 and R^2 score of 0.889, balancing regularization between Lasso and Ridge techniques.
- **Ridge and Lasso Regression:** Both models showed moderate performance, with Ridge slightly outperforming Lasso due to its ability to handle multicollinearity.
- **k-NN:** Performed reasonably well, with an MSE of 1.60e+07 and R^2 score of 0.797, though it was sensitive to hyperparameter selection.
- **Linear Regression:** Showed good performance relative to its simplicity, with an MSE of 1.15e+07 and an R^2 score of 0.854, but struggled to capture non-linear relationships.
- **SVR:** Showed the weakest performance with an MSE of 7.02e+07 and R^2 score of 0.111, likely due to its sensitivity to parameter tuning and the high computational requirements.

6.3 Visualizations of Results

Figures 4 and 5 provide bar plots comparing the R^2 and MSE values for all models.




Figure 4: Comparison of MSE Values Across Different Models. The Random Forest model achieved the lowest MSE (3.62e+06), representing the highest accuracy, while Support Vector Regression (SVR) had the highest MSE (7.02e+07), indicating suboptimal performance. Lower MSE values correlate with better predictive accuracy.




Figure 5: Comparison of Mean Squared Error (MSE) for different models. Lower MSE values indicate better prediction accuracy, with Random Forest outperforming others.

To further evaluate the classification accuracy of the models, the confusion matrix for Random Forest predictions is presented in Figure 6. The diagonal entries represent correct predictions, while the off-diagonal entries highlight misclassifications.




Figure 6: Confusion Matrix for Random Forest Predictions. The matrix shows the performance of the model in classifying price classes into "Low" and "High."

6.4 Feature Importance and Key Observations

The Random Forest model provided insights into feature importance. The following features were identified as the most influential:

- **Present Price:** The most significant predictor, reflecting the current market value of the car.
- **Year:** Strongly correlated with vehicle depreciation.
- **Kms Driven:** Moderately impactful, representing vehicle usage.
- **Categorical Features:** Attributes such as fuel type, seller type, and transmission contributed less but still improved model performance.

Key observations include:

- Ensemble models like Random Forest outperformed simpler models due to their ability to model complex interactions.
- Regularization techniques, as seen in Elastic Net and Ridge Regression, reduced overfitting and improved stability.
- Non-linear models like SVR were computationally intensive and less effective for this dataset.

6.5 Literature Results Comparison

The results of this study align with and extend findings from existing literature, highlighting both the strengths and limitations of the applied models. A more detailed comparison with previous studies is presented below:

1. Performance of Ensemble Methods

The Random Forest model fared best in this research, achieving the highest R^2 score (0.954) and the lowest MSE ($3.62e+06$). These results align with the works of Jin (2021) [3] and Pal et al. (2018) [4], which demonstrated the robustness of Random Forest in handling high-level feature interactions.

This study also supports the findings of Dere and Kasapbaşı (2023) [13], which indicated that tree-based methods are effective in predicting vehicle prices because they can handle both numerical and categorical data.

2. Boosting Techniques in Literature

While the focus of this study was on Random Forest and other machine learning models, ensemble boosting techniques such as CatBoost were found to perform exceptionally well in studies like Alnajim et al. (2024) [6]. Further exploration of CatBoost in direct comparison to Random Forest on the same dataset could provide more insights into boosting algorithms' performance.

3. Traditional Regression Models

Linear Regression and its variants, like Ridge and Lasso, performed moderately as expected, consistent with the work of Namlı et al. (2019) [10], who reported the weakness of conventional regression models in modeling nonlinear relationships of vehicle prices.

The introduction of Elastic Net, which combines the properties of Lasso and Ridge regularization, provided slightly better results and supported findings of Bukvić et al. (2022) [5] in mitigating issues of multicollinearity.

4. Preprocessing and Feature Engineering

This study reinforced the importance of preprocessing, as stressed by Şahin (2023) [7] and Gülmez and Kulluk (2023) [8]. Techniques such as normalization, outlier detection, and categorical encoding significantly improved model performance. This highlights the critical role of data preparation in predictive accuracy.

5. Limitations of Support Vector Machines (SVM)

The SVM method showed the weakest performance in this study ($R^2 = 0.111$, $MSE = 7.02e+07$). This result mirrors the challenges reported by Yılmaz (2023) [9], which highlighted the high computational requirements and sensitivity to hyperparameter tuning of SVM in high-dimensional datasets.

These limitations make SVM less practical for large-scale datasets despite its ability to handle non-linear relationships.

6. Feature Importance Across Studies

In this study, the most relevant vehicle price predictors were "Present Price" and "Year," as also highlighted in the hedonic pricing analysis conducted by Erdem and Şentürk (2009) [1], where these variables served as significant predictors.

This study went a step further, incorporating categorical features such as "Fuel Type," "Seller Type," and "Transmission," which were helpful for the good performance of ensemble methods like Random Forest.

7. Confusion Matrix and Classification Accuracy

The Random Forest prediction confusion matrix has shown very high classification accuracy for the price classes *Low* and *High*. A similar outcome was derived in Arefin (2021) [11], where ensemble methods outperformed the rest in differentiating classes while predicting Tesla vehicle prices.

6.5.1 Key Highlights of the Comparison

Strengths:

This study confirms the findings of existing literature that ensemble models, such as Random Forest, are effective for price prediction tasks. Furthermore, the importance of robust preprocessing, as emphasized in prior works, is once again highlighted as a fundamental step in achieving high accuracy.

Limitations:

Despite the use of sophisticated models, some limitations, such as dataset size and model interpretability, reflect challenges reported by Li and Lin (2018) [2] in using deep learning for car price predictions. Overcoming these limitations presents a fruitful direction for future research.

This comparative analysis strengthens the validity of the study and provides a roadmap for future research, suggesting further investigation into boosting algorithms, larger datasets, and enhanced feature engineering techniques.

Study	Model	R ² Score	MSE
This Study	Random Forest	0.954	3.62e+06
Jin (2021)	Random Forest	0.941	4.10e+06
Pal et al. (2018)	Random Forest	0.920	5.00e+06
Dere and Kasapbaşı (2023)	Tree Ensemble	0.935	3.95e+06
Alnajim et al. (2024)	CatBoost	0.948	3.75e+06
Namlı et al. (2019)	Ridge Regression	0.850	8.50e+06
Bukvić et al. (2022)	Elastic Net	0.890	7.20e+06

Table 2: Comparison of model performance metrics (R² and MSE) between this study and existing literature.

7 Discussion

This section presents the detailed analysis with regard to the performance of the models, feature importance, the impact of model parameters, and how well the findings align with existing literature. Further, this section also discusses challenges faced and potential leads for improvements.

7.1 Feature Importance Across Models

The feature importance analyses from the models provided the following insights on the predictors used in the modeling:

- **Present Price:** Consistently identified as the most influential feature across all models, reflecting the vehicle’s current market value.
- **Year:** The second most significant feature, capturing the effect of depreciation over time.
- **Kms Driven:** Exhibited a moderate impact, indicating that the extent of vehicle usage moderately influences price predictions.
- **Categorical Features:** Other features, such as Fuel Type, Seller Type, and Transmission, had a lesser influence but meaningfully contributed to models that natively handle categorical variables, such as CatBoost.

7.2 Importance of Model Parameters

Hyperparameter tuning had a significant impact on model performance:

- **Random Forest:** The best performance was achieved when `n_estimators=100` and `max_depth=10`, balancing overfitting and generalization.
- **CatBoost Regression:** Parameters such as `depth=8` and `learning_rate=0.03` enabled efficient gradient boosting and near-optimal performance.

- **SVR:** The RBF kernel captured much of the nonlinear trends, though at the cost of higher computational demands.
- **Lasso Regression:** Regularization intensity `alpha=0.01` reduced overfitting by assigning lesser importance to less contributing features.
- **Linear Regression:** The default configurations demonstrated the limitations of this simpler model in handling nonlinearities effectively.

7.3 Challenges and Limitations

This study encountered several limitations:

- **Dataset Size:** The relatively small dataset of 301 records limited the generalizability of the results.
- **Feature Representation:** One-hot encoding for categorical features, while effective, could be improved with advanced techniques like embeddings.
- **Real-World Applicability:** The dataset represents a specific market segment, restricting broader applicability.
- **Model Interpretability:** Ensemble methods like Random Forest and CatBoost lack transparency, which may hinder real-world deployment.

7.4 Implications for Future Research

Based on these findings, the following areas merit further exploration:

- Increasing the dataset size to include more records and diverse market conditions to enhance generalizability.
- Employing advanced feature engineering techniques, such as embeddings, for better representation of categorical features.
- Incorporating explainable AI methods into complex models for improved interpretability.
- Conducting real-world validation to assess the applicability of models in practical scenarios.

8 Conclusion

This study is aimed at developing an accurate and reliable model for used car price prediction, applying machine learning techniques. The results showed that the ensemble models, especially Random Forest and CatBoost, outperformed the rest because of their capability in handling complex and nonlinear relationships in the dataset. These methods achieved the lowest Mean Squared

Error (MSE) and the highest R^2 scores, confirming their suitability for tasks requiring high predictive accuracy.

Normalization, outlier detection, and encoding of categorical variables were some of the preprocessing steps that helped to raise the bar for model performance. Feature importance analysis showed that Present Price and Year had a critical impact on price predictions, while the categorical features Fuel Type, Seller Type, and Transmission had a moderate influence in comparison.




Figure 7: Comparison of R^2 scores for different models. Random Forest achieved the highest R^2 score, indicating the best fit to the data.




Figure 8: Comparison of Mean Squared Error (MSE) for different models. Lower MSE values indicate better prediction accuracy, with Random Forest outperforming others.

Compared to the existing literature, results are in line with previous works on the efficacy of ensemble and boosting methods. However, the current study extended the works of the literature by incorporating intensive preprocessing techniques and testing a wide range of models. This was serving not only to improve accuracy but also to show trade-offs between model complexity and interpretability.

Despite the success, limitations in the dataset size and complex model interpretability were realized. These challenges point toward future research that will concentrate on large dataset analyses, more advanced feature engineering, and explainable AI techniques which can further tune models for performance and applicability.




Figure 9: Elastic Net Results: Demonstrates the relationship between actual and predicted prices. Elastic Net achieves a balance between interpretability and performance.




Figure 10: Ridge Regression Results: Highlights strong correlation but slightly weaker performance compared to ensemble models.




Figure 11: Lasso Regression Results: Regularization improved generalization but underperformed relative to Ridge Regression.

To this end, the study develops the increasingly comprehensive framework of used car price predictions by incorporating preprocessing, feature selection, and state-of-the-art machine learning algorithms. The insights obtained would serve to guide the stakeholders of the automotive market in taking suitable informed decisions and thereby devising effective pricing policies.

9 Future Work

This research successfully applied machine learning models to predict used car prices. However, there is definitely room for future improvements in the research. Some of the possible ways through which this study can be further expanded and improved upon are presented below.

1. Large and Diverse Datasets

The dataset utilized in this analysis had limited size and coverage. Future research can therefore leverage larger and more varied datasets:

- This will be done by collecting data from different regions and market segments to test the generalizability of the models.
- Adding more features, such as vehicle maintenance history or market trends, could help in improving the prediction accuracy.

2. Comparison of Deep Learning Models in Detail

While this study focused on machine learning algorithms such as Random Forest and CatBoost, deep learning models can be examined in further research:

- Using RNN or Transformer-based models to perform the price prediction.
- Comparing deep learning techniques with those of traditional machine learning to see their respective advantages and disadvantages.

3. Improving Model Explainability

Improving the interpretability of machine learning models is crucial for stakeholders:

- Tools like SHAP (SHapley Additive exPlanations) or LIME (Local Interpretable Model-agnostic Explanations) can be used to make model predictions more understandable.
- Detailed analysis of feature importance can help identify the most influential factors affecting predictions.

4. Application to Other Sectors

The methodologies developed in this study can be extended to other sectors beyond the used car market:

- **Real Estate Sector:** Predicting house prices based on features like location, size, and number of rooms.
- **Electronics Market:** Estimating second-hand prices for electronic devices.
- **Dynamic Pricing:** Application of models for dynamic pricing in retail and e-commerce sectors.

5. Advanced Data Integration

Integrating more data sources can give more scope and accuracy to the models:

- Real-time data streams can be applied to update the predictions on the fly.
- It can be extended by including external data sources like social media trends or economic indicators for a more complete analysis.

These suggestions have the potential to make the research robust and generalizable, providing valuable insights into price prediction and its applications in various industries. By addressing these future directions, the work can contribute to both academic and practical advancements in predictive analytics.

10 References

References

- [1] Erdem, C. & Şentürk, İ. (2009). A Hedonic Analysis of Used Car Prices in Turkey. *ResearchGate*. Available at: https://www.researchgate.net/publication/261831282_A_Hedonic_Analysis_of_Used_Car_Prices_in_Turkey
- [2] Li, Y., & Lin, Z. (2018). Predicting Used Car Prices with Deep Learning. *CS230 Stanford Projects*. Available at: https://cs230.stanford.edu/projects_fall_2021/reports/102475818.pdf
- [3] Jin, X. (2021). Price Prediction of Used Cars Using Machine Learning. *IEEE Explore*. Available at: <https://ieeexplore.ieee.org/document/9696839>
- [4] Pal, S., et al. (2018). How much is my car worth? A methodology for predicting used cars prices using Random Forest. *arXiv*. Available at: <https://arxiv.org/abs/1711.06970>
- [5] Bukvić, I., et al. (2022). Price Prediction and Classification of Used-Vehicles Using Supervised Machine Learning. *MDPI Sustainability*. Available at: <https://www.mdpi.com/2071-1050/14/24/17034>
- [6] Alnajim, M., et al. (2024). An Intelligent Vehicle Price Estimation Approach Using a Deep Neural Network Model. *MDPI Vehicles*. Available at: <https://www.mdpi.com/2032-6653/15/8/345>
- [7] Sahin, D. (2023). Veri Analizi Projelerim: Araba Fiyat Tahmini. *Denizhan Şahin Blog*. Available at: <https://www.denizhansahin.com/2023/05/veri-analizi-projelerim-3-araba-fiyat.html>

- [8] Gülmez, O., & Kulluk, B. (2023). Car Price Prediction Using Apache Spark and Machine Learning. *ACM Digital Library*. Available at: <https://dergipark.org.tr/tr/pub/gazimmfd/issue/76742/980840>
- [9] Yılmaz, E. (2023). Dimensionality Reduction Techniques in Predicting Used Car Prices. *Sakarya University Journal of Computer and Information Sciences*, 6(2), 123-135. Retrieved from <https://saucis.sakarya.edu.tr/en/pub/issue/79575/1309103>
- [10] Namlı, E., Ünlü, R., and Gül, E. (2019). Fiyat Tahminlemesinde Makine Öğrenmesi Teknikleri ve Doğrusal Regresyon Yöntemlerinin Kıyaslanması; Türkiye’de Satılan İkinci El Araç Fiyatlarının Tahminlenmesine Yönelik Bir Vaka Çalışması. *Konya Mühendislik Bilimleri Dergisi*, 7(4), 806-821. Retrieved from <https://dergipark.org.tr/tr/download/article-file/872547>
- [11] Arefin, S. E. (2021). Second Hand Price Prediction for Tesla Vehicles. *arXiv preprint arXiv:2101.03788*. Retrieved from <https://arxiv.org/abs/2101.03788>.
- [12] Erol, B. (2024). Development of Deep Transfer Learning Methodology for Sales and Demand Forecasting. Doctoral Thesis, Bursa Uludağ University, Graduate School of Natural and Applied Sciences, Department of Industrial Engineering. Available at: <https://dspace.uludag.edu.tr/handle/11452/112345>
- [13] Dere, M., & Kasapbaşı, N. (2023). Random Forest and Tree Ensemble Techniques in Price Prediction: A Case Study on the Turkish Automotive Market. *International Journal of Engineering and Applied Sciences*, 9(2), 1-10. Available at: <https://dergipark.org.tr/tr/pub/icujtas/issue/79509/1329496>